

Logit-lens traces: $\text{logit}(\text{original}) - \text{logit}(\text{mutated})$ per transformer layer
 Positive = model prefers training-time name; negative = model follows documentation

